

Reinforcement Learning

CS 59300: RL1

September 4, 2025

Joseph Campbell
Department of Computer Science

Today's lecture


1. Evaluating Policy Quality
2. Known Models: Planning with Dynamic Programming
3. Unknown Models: Monte Carlo Learning

Some content inspired by David Silver's UCL RL course and Katerina Fragkiadaki's CMU 10-403

Recap

The agent and environment operate at discrete timesteps $t = 0, 1, 2, \dots$

- The agent observes state s_t at time t
- The agent takes action a_t
- The agent gets the resulting reward r_{t+1} and the subsequent state s_{t+1}




Recap

Action a_t is chosen by sampling actions from a probability distribution

$$a_t \sim \pi(a|s)$$

The probability distribution π is referred to as a **policy**.



Recap: Markov Decision Process

We formulate sequential decision-making problems as MDPs

Definition: a tuple consisting of $(\mathcal{S}, \mathcal{A}, R, T, \gamma)$:

- \mathcal{S} is the set of states in our environment
- \mathcal{A} is a set of actions that can be taken by an agent
- R is the reward function. $\mathcal{S} \times \mathcal{A} \mapsto \mathbb{R}$
- T is the state transition probability. $\mathcal{S} \times \mathcal{S} \times \mathcal{A} \mapsto [0,1]$
- γ is a discount factor. $\gamma \in [0,1]$

Value functions

Definition: the **action-value** function of an MDP is the expected return starting from state s , taking action a , and following policy π for all subsequent states

$$Q_{\pi}(s, a) = \mathbb{E}[G_t | s_t = s, a_t = a]$$

Definition: the **state-value function** of an MDP is the expected return starting from state s and following policy π

$$V_{\pi}(s) = \mathbb{E}[G_t | s_t = s]$$

Recap: Finding an optimal policy

If we know the optimal action-value function $Q^*(s, a)$...

$$\pi^*(a|s) = \begin{cases} 1, & \text{if } a = \operatorname{argmax}_{a \in \mathcal{A}} Q^*(s, a) \\ 0, & \text{otherwise} \end{cases}$$

...we immediately have the optimal policy. We simply follow it.

Recap: Finding an optimal policy

If we know the optimal state-value function $V^*(s)$...

$$\pi^*(a|s) = \begin{cases} 1, & \text{if } a = \operatorname{argmax}_{a \in \mathcal{A}} \left[\sum_{s', r} p(s', r | s, a) (r + \gamma V^*(s')) \right] \\ 0, & \text{otherwise} \end{cases}$$

...we need access to the state transition function (defined in the MDP earlier as T) to do one-step ahead lookup

- Take the action which leads us to the state with highest value

Evaluating Policy Quality

The reward hypothesis

Let's take a step back: **why do we care about rewards and values?**

The reward hypothesis

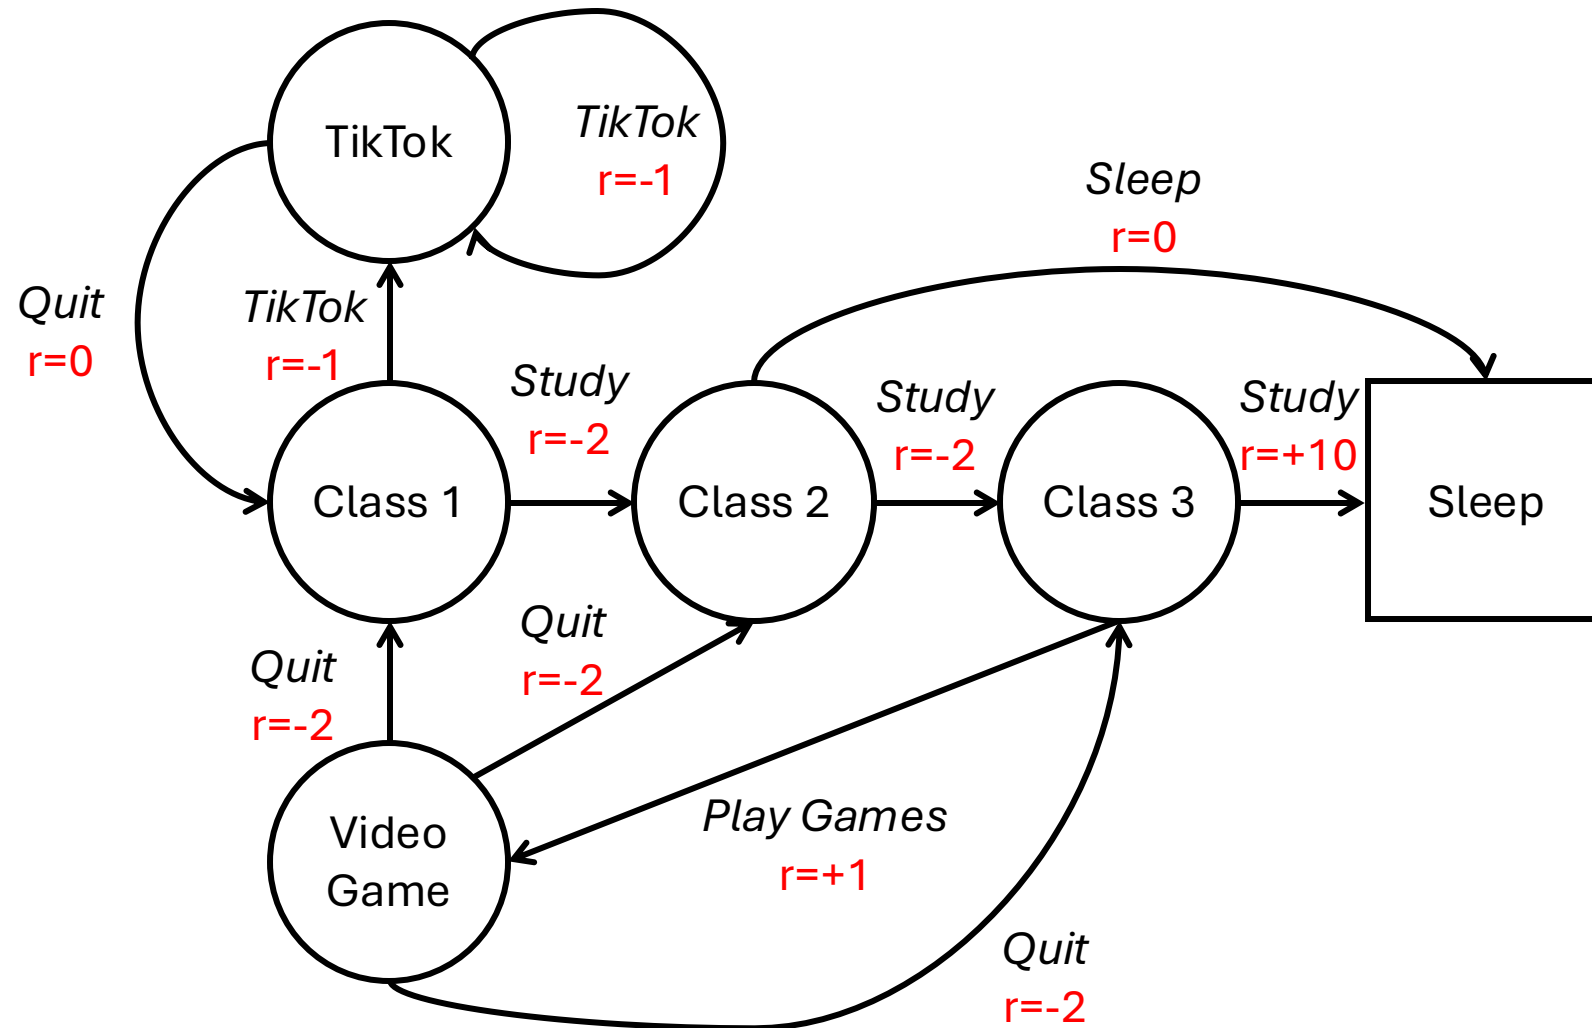
Let's take a step back: **why do we care about rewards and values?**

“That all of what we mean by goals and purposes can be well thought of as the maximization of the expected value of the cumulative sum of a received scalar signal (called reward).”

- Sutton and Barto, Chapter 3.2

Do you agree? Why or why not?

Example: student life as an MDP



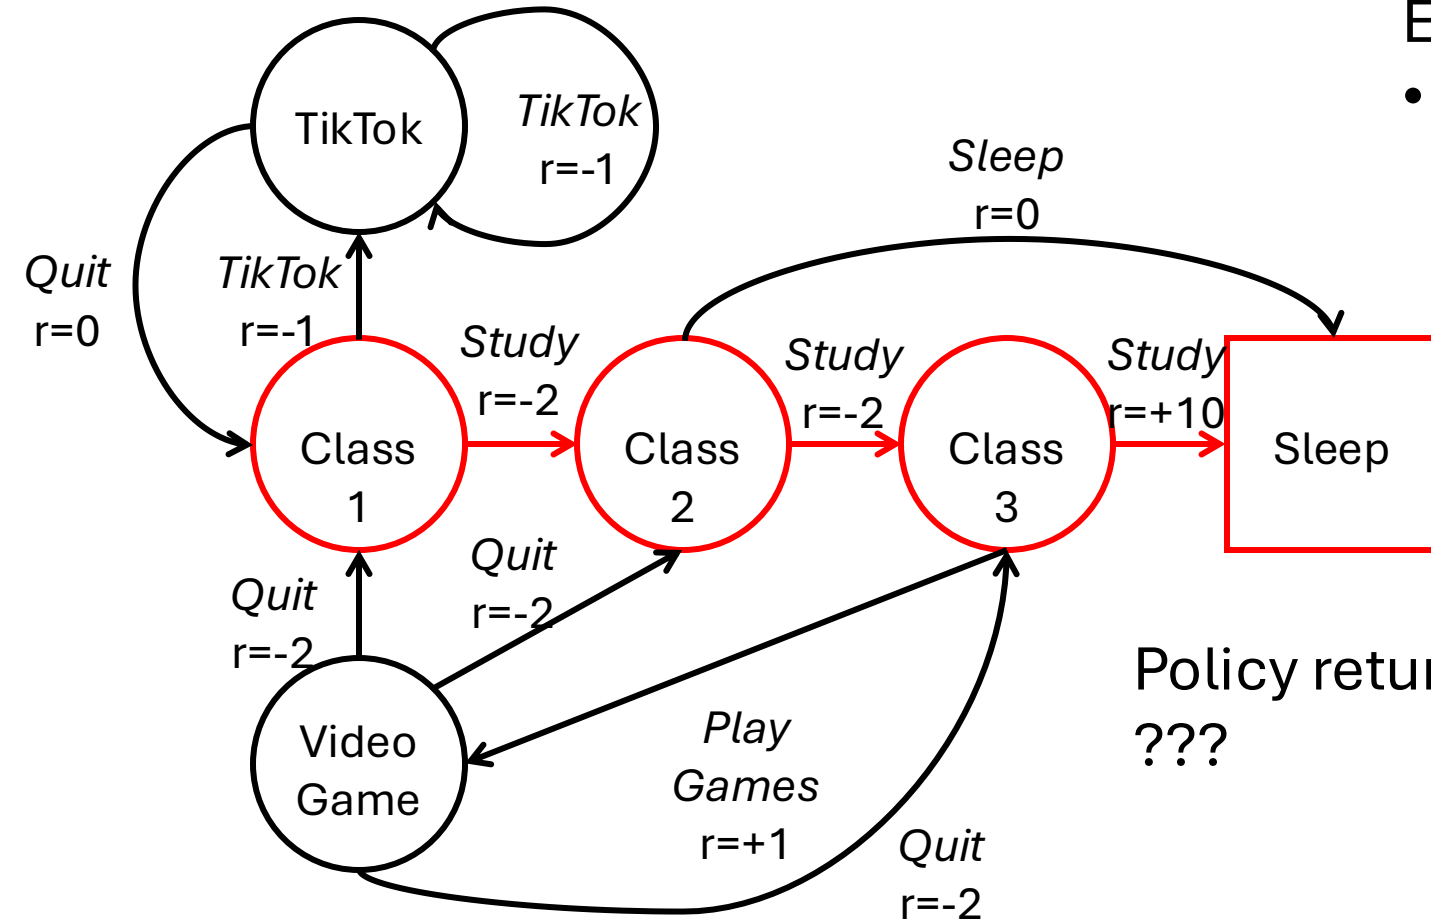
Recall: Returns for a policy

$$G_t = r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \gamma^3 r_{t+4} + \dots$$

Example: student life as an MDP

Example policy rollout:

- Class 1, Class 2, Class 3, Sleep

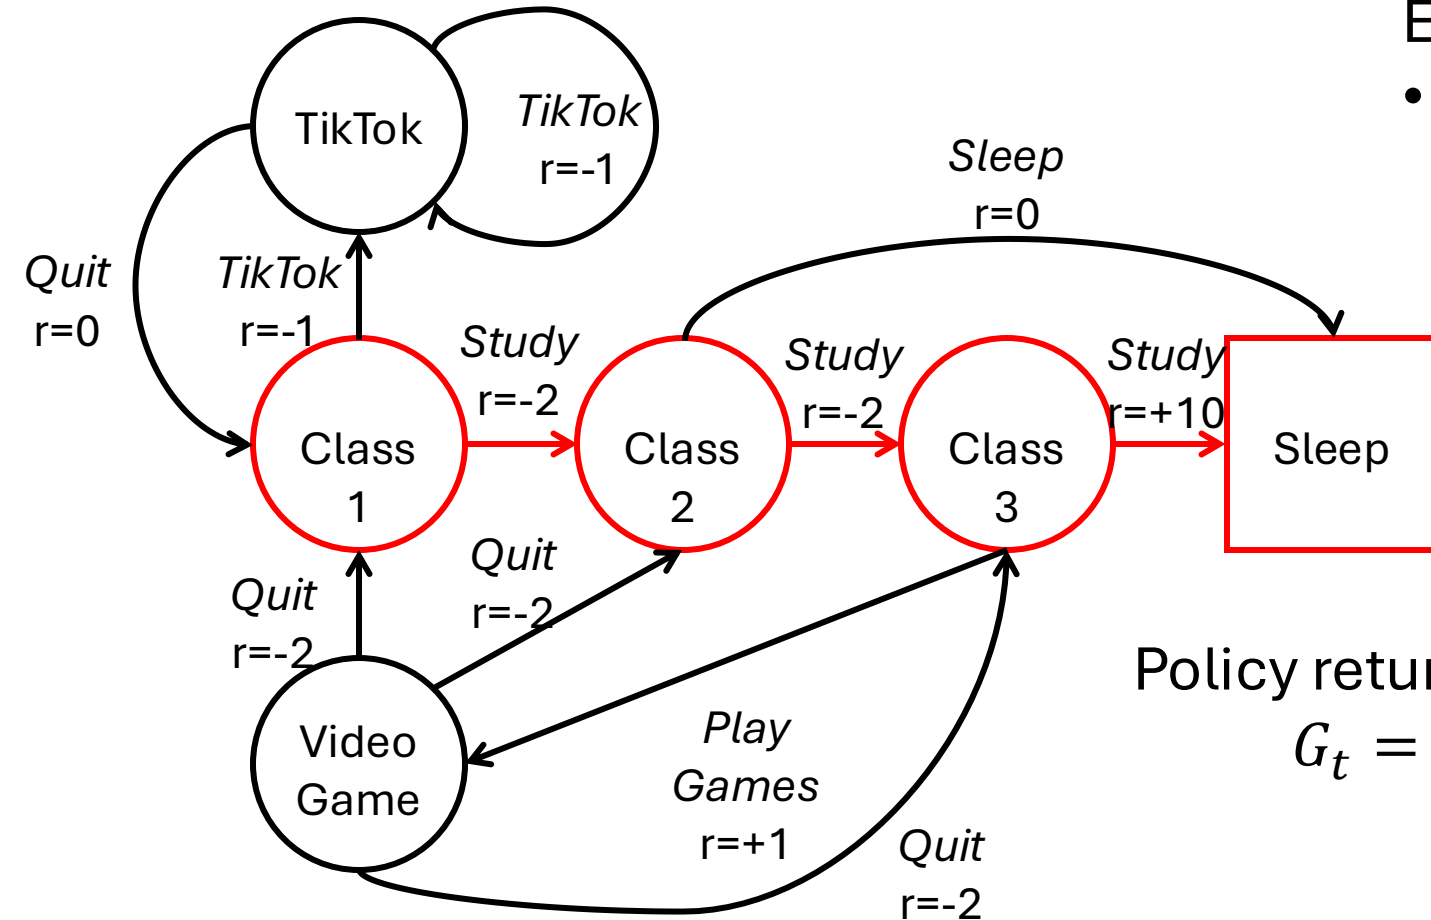


Policy return:
???

Example: student life as an MDP

Example policy rollout:

- Class 1, Class 2, Class 3, Sleep



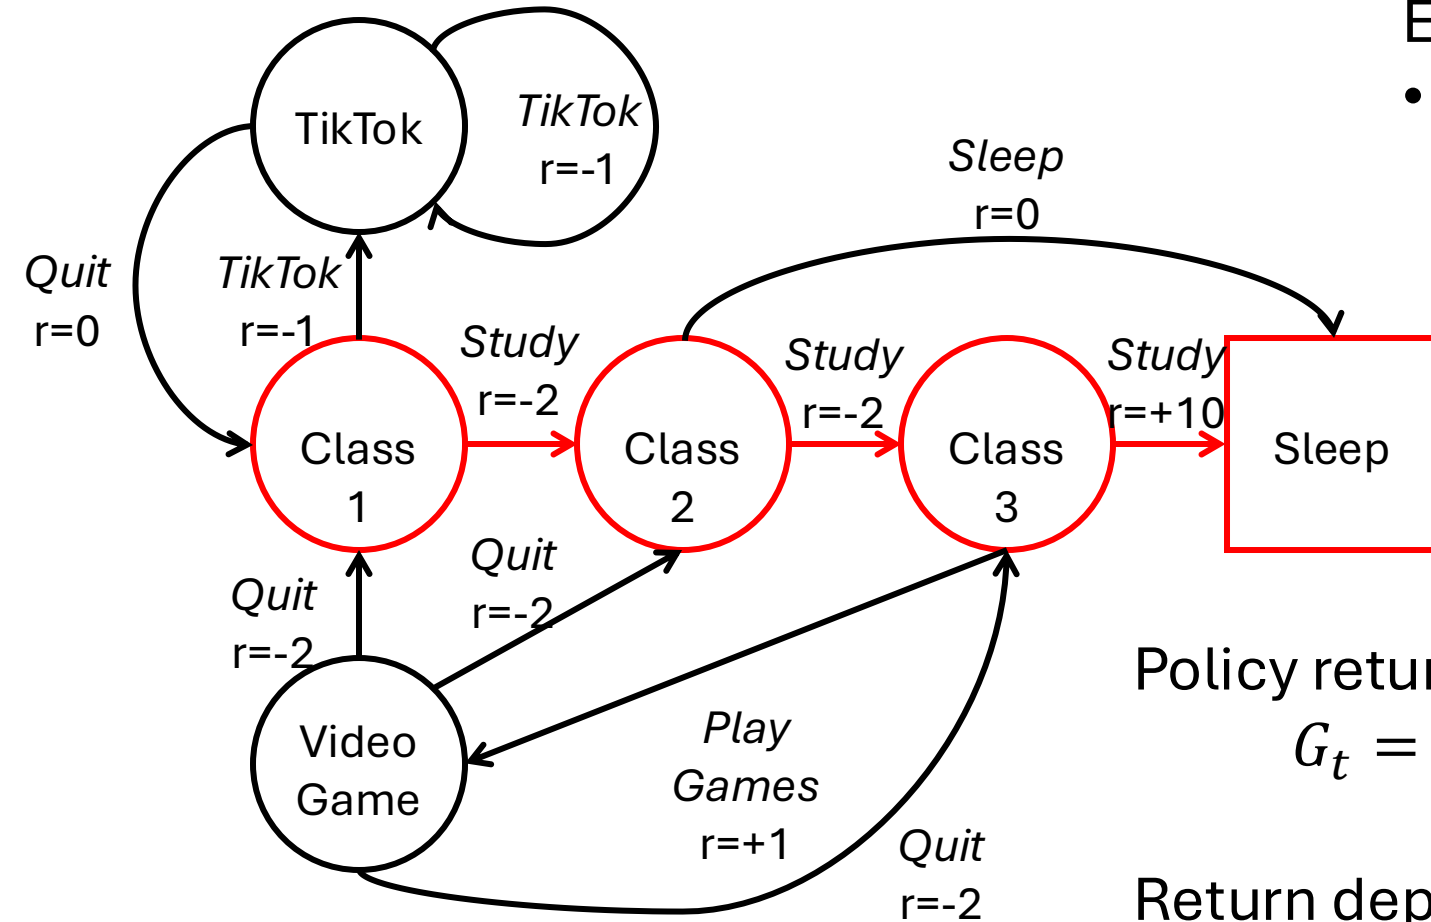
Policy return:

$$G_t = -2 + \gamma(-2) + \gamma^2(+10)$$

Example: student life as an MDP

Example policy rollout:

- Class 1, Class 2, Class 3, Sleep



Policy return:

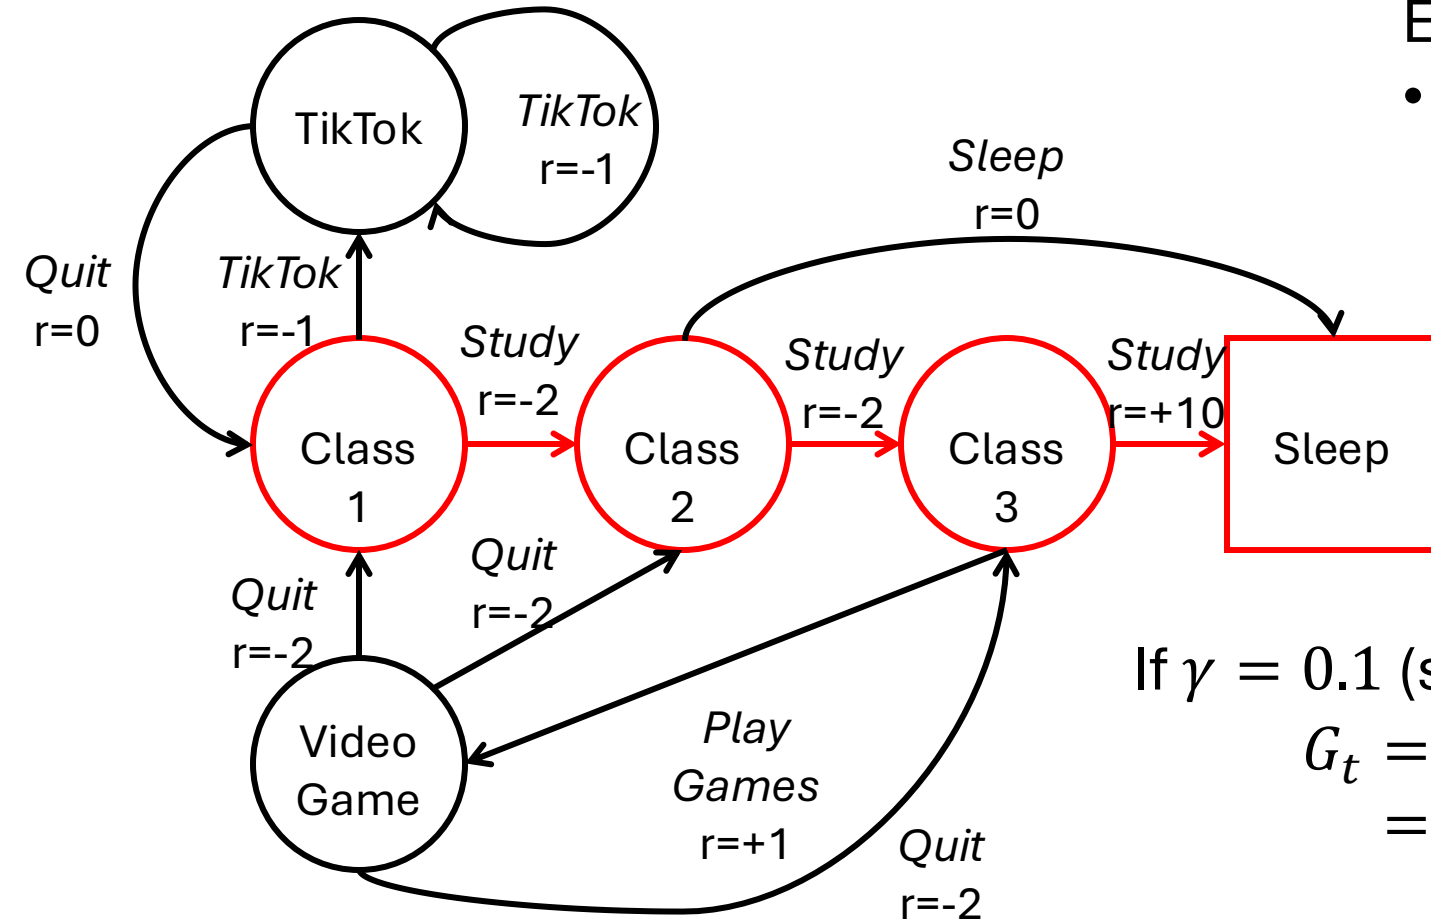
$$G_t = -2 + \gamma(-2) + \gamma^2(+10)$$

Return depends on γ !

Example: student life as an MDP

Example policy rollout:

- Class 1, Class 2, Class 3, Sleep



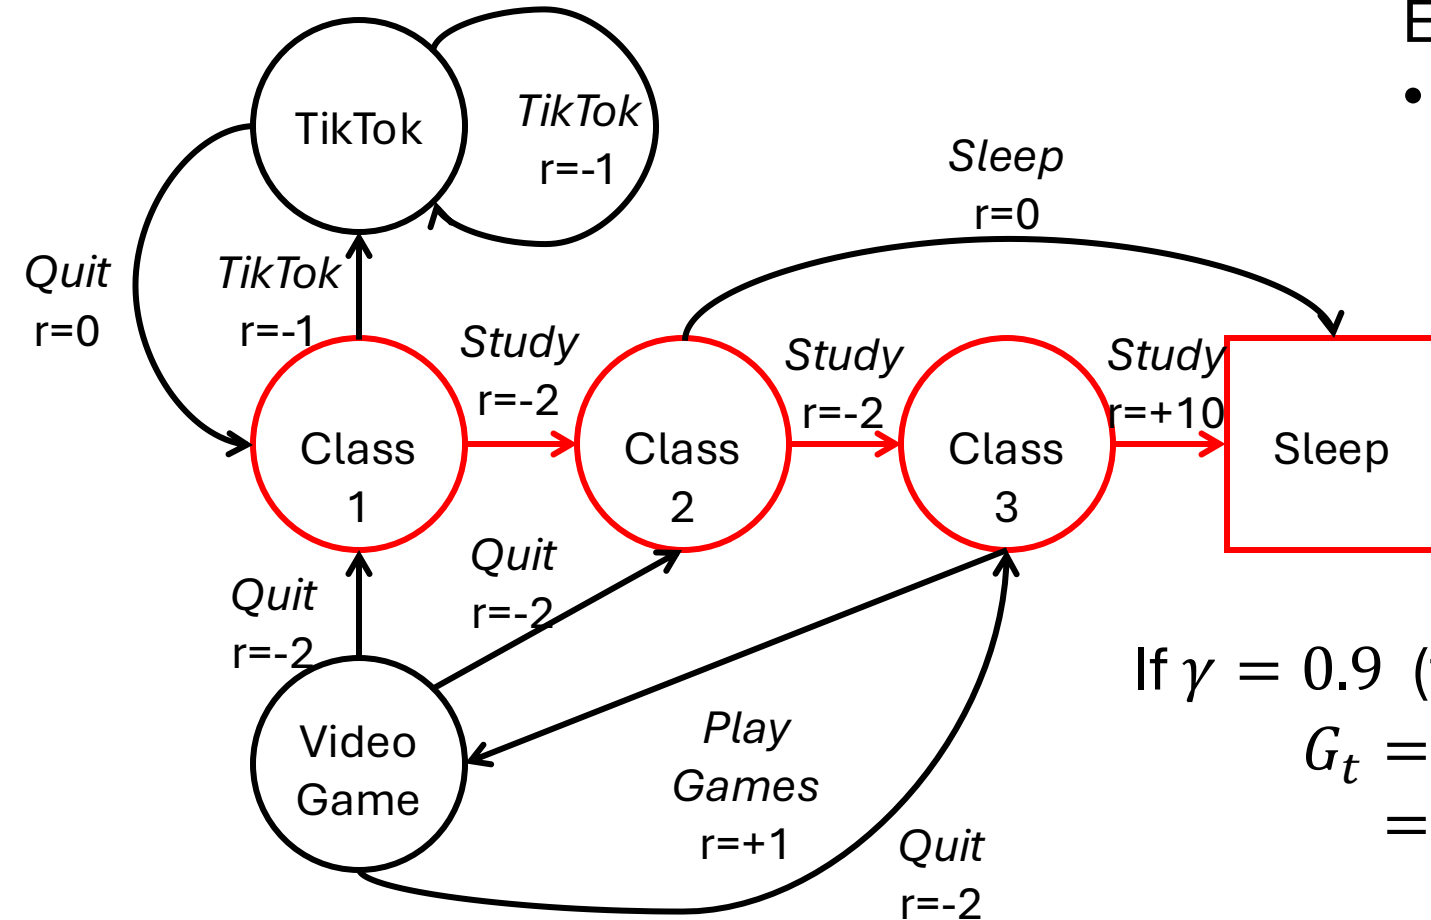
If $\gamma = 0.1$ (short-sighted):

$$\begin{aligned} G_t &= -2 + 0.1(-2) + 0.1^2(+10) \\ &= -2.1 \end{aligned}$$

Example: student life as an MDP

Example policy rollout:

- Class 1, Class 2, Class 3, Sleep



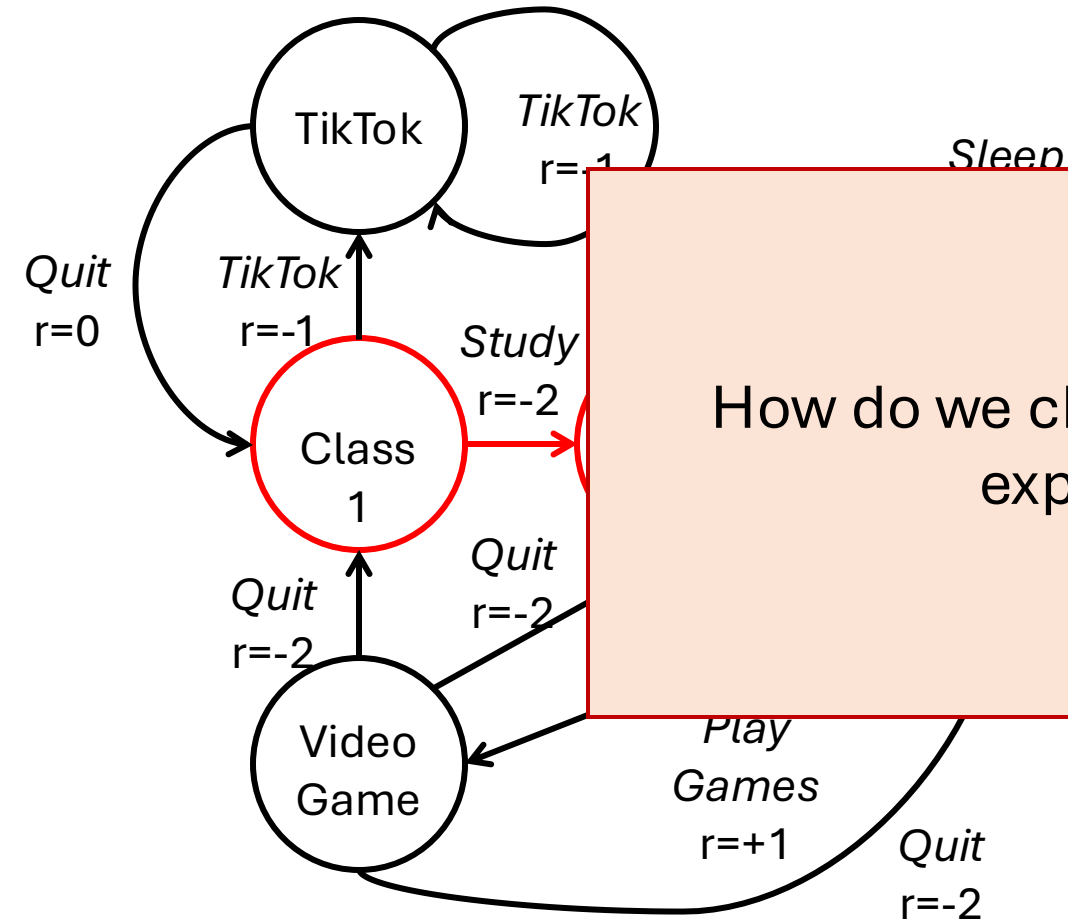
If $\gamma = 0.9$ (far-sighted):

$$\begin{aligned} G_t &= -2 + 0.9(-2) + 0.9^2(+10) \\ &= +4.3 \end{aligned}$$

Example: student life as an MDP

Example policy rollout:

- Class 1, Class 2, Class 3, Sleep



How do we choose actions so as to maximize the expected return for a given γ ?

$$\begin{aligned} G_t &= -2 + 0.9(-2) + 0.9^2(-2) + 0.9^3(+10) \\ &= +1.87 \end{aligned}$$

Recursive form of policy returns

$$G_t = r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \gamma^3 r_{t+4} + \dots$$

Recursive form of policy returns

$$\begin{aligned} G_t &= r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \gamma^3 r_{t+4} + \dots \\ &= r_{t+1} + \gamma(r_{t+2} + \gamma r_{t+3} + \gamma^2 r_{t+4} + \dots) \end{aligned}$$

Recursive form of policy returns

$$\begin{aligned} G_t &= r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \gamma^3 r_{t+4} + \dots \\ &= r_{t+1} + \gamma(r_{t+2} + \gamma r_{t+3} + \gamma^2 r_{t+4} + \dots) \\ &= r_{t+1} + \gamma G_{t+1} \end{aligned}$$

Recursive form of policy returns

$$\begin{aligned} G_t &= r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \gamma^3 r_{t+4} + \dots \\ &= r_{t+1} + \gamma(r_{t+2} + \gamma r_{t+3} + \gamma^2 r_{t+4} + \dots) \\ &= r_{t+1} + \gamma G_{t+1} \end{aligned}$$

This relationship allows us to decompose value functions

Bellman expectation equation

We can decompose value functions into two parts:

- The immediate reward
- The expected future returns

$$\mathbb{E}[G_t | s_t = s] = \mathbb{E}[r_{t+1} + \gamma G_{t+1} | s_t = s]$$

Bellman expectation equation

We can decompose value functions into two parts:

- The immediate reward
- The expected future returns

$$\mathbb{E}[G_t | s_t = s] = \mathbb{E}[r_{t+1} + \gamma G_{t+1} | s_t = s]$$

State-value: $V_{\pi}(s) = \mathbb{E}[r_{t+1} + \gamma V_{\pi}(s_{t+1}) | s_t = s]$

Bellman expectation equation

We can decompose value functions into two parts:


- The immediate reward
- The expected future returns

$$\mathbb{E}[G_t | s_t = s] = \mathbb{E}[r_{t+1} + \gamma G_{t+1} | s_t = s]$$

State-value: $V_\pi(s) = \mathbb{E}[r_{t+1} + \gamma V_\pi(s_{t+1}) | s_t = s]$


Action-value: $Q_\pi(s, a) = \mathbb{E}[r_{t+1} + \gamma Q_\pi(s_{t+1}, a_{t+1}) | s_t = s, a_t = a]$

State-value and action-values are related




$$V_\pi(s) = \sum_{a \in \mathcal{A}} \pi(a|s) Q_\pi(s, a)$$

State-value and action-values are related




$$Q_\pi(s, a) = r + \gamma \sum_{s' \in \mathcal{S}} p(s'|s, a) V_\pi(s')$$

Bellman expectation equation for state-value



$$V_\pi(s) = \sum_{a \in \mathcal{A}} \pi(a|s) \left(r + \gamma \sum_{s' \in \mathcal{S}} p(s'|s, a) V_\pi(s') \right)$$

Bellman expectation equation for action-value

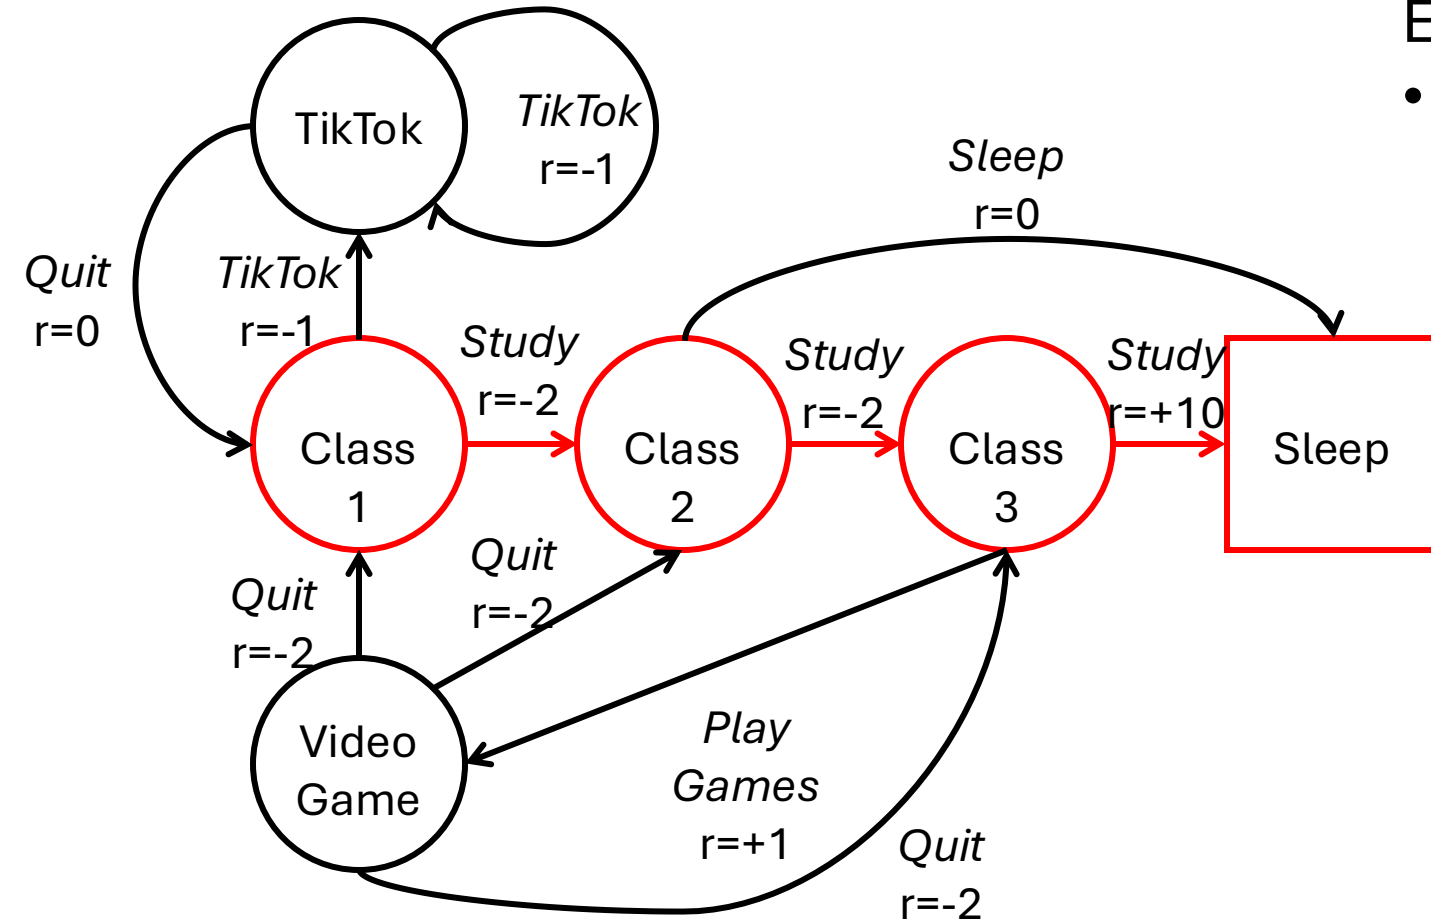


$$Q_\pi(s, a) = r + \gamma \sum_{s' \in \mathcal{S}} p(s'|s, a) \sum_{a' \in \mathcal{A}} \pi(a'|s') Q_\pi(s', a')$$

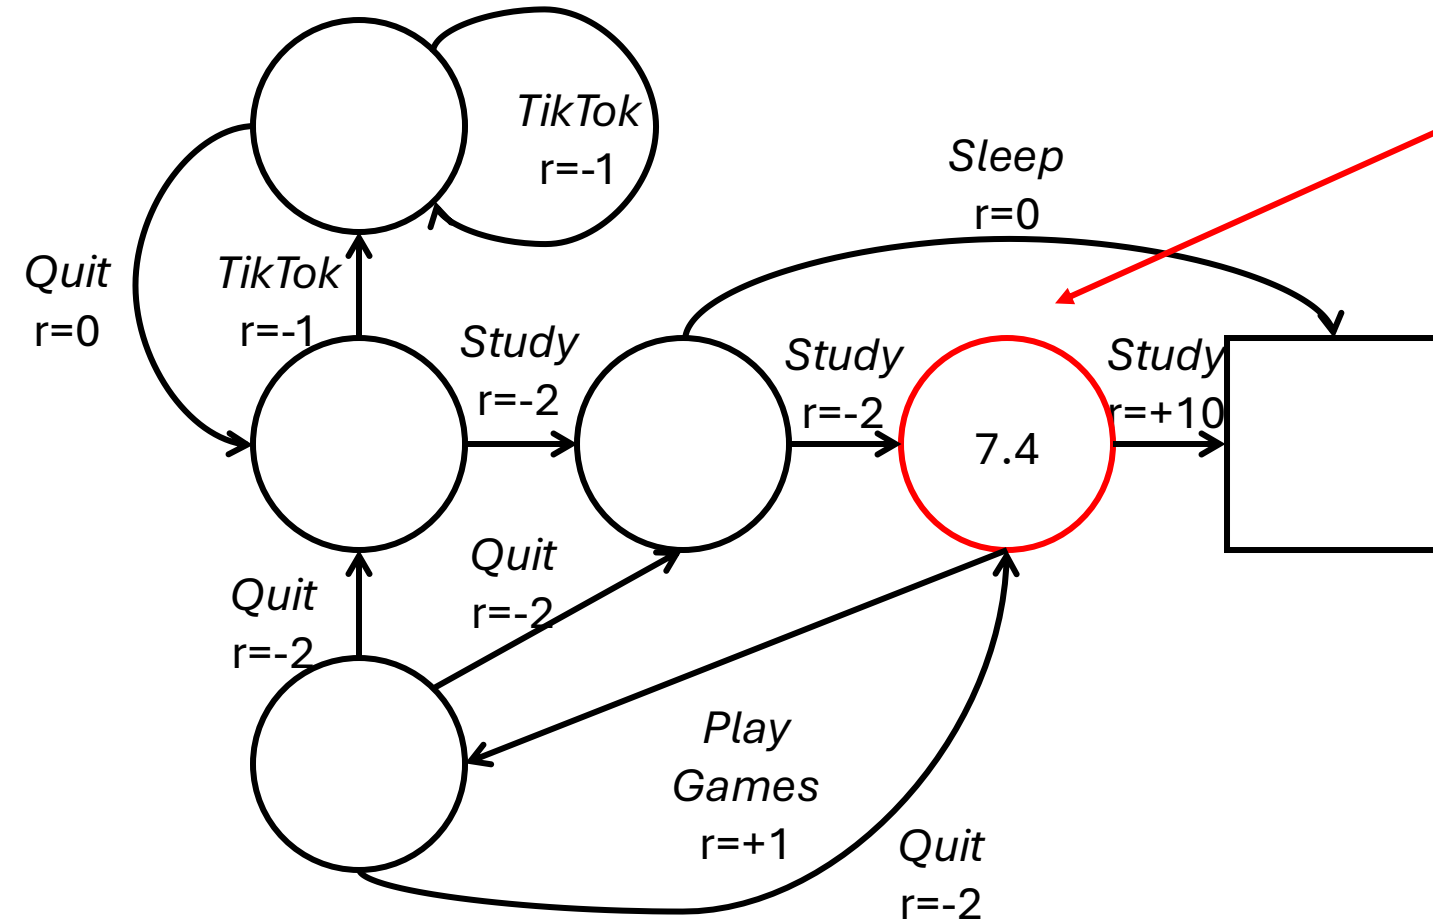
Evaluating a policy with Bellman expectation

Example policy rollout:

- Class 1, Class 2, Class 3, Sleep

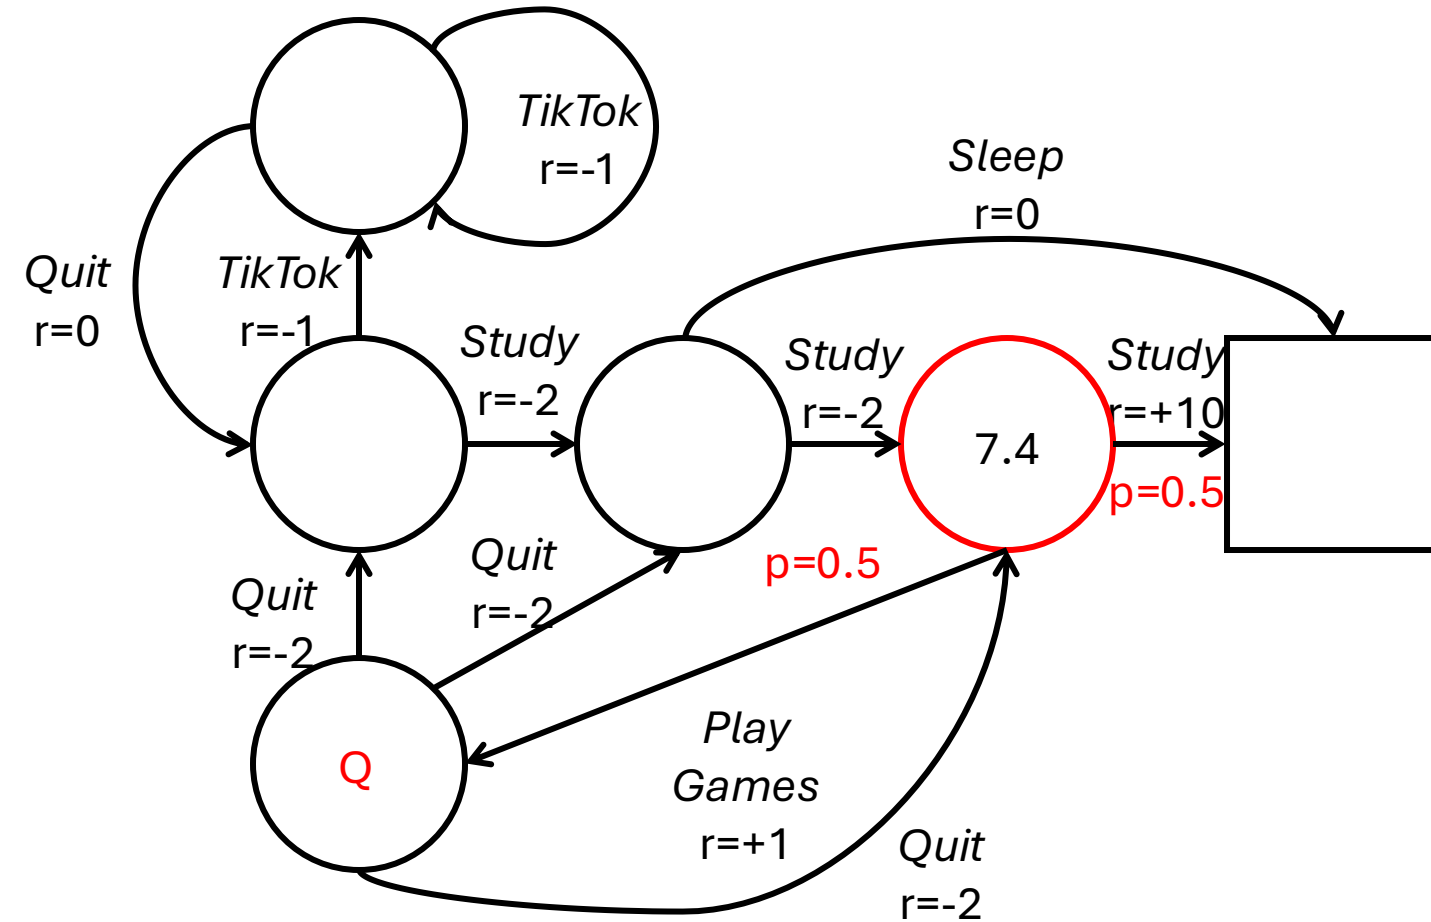


Evaluating a policy with Bellman expectation



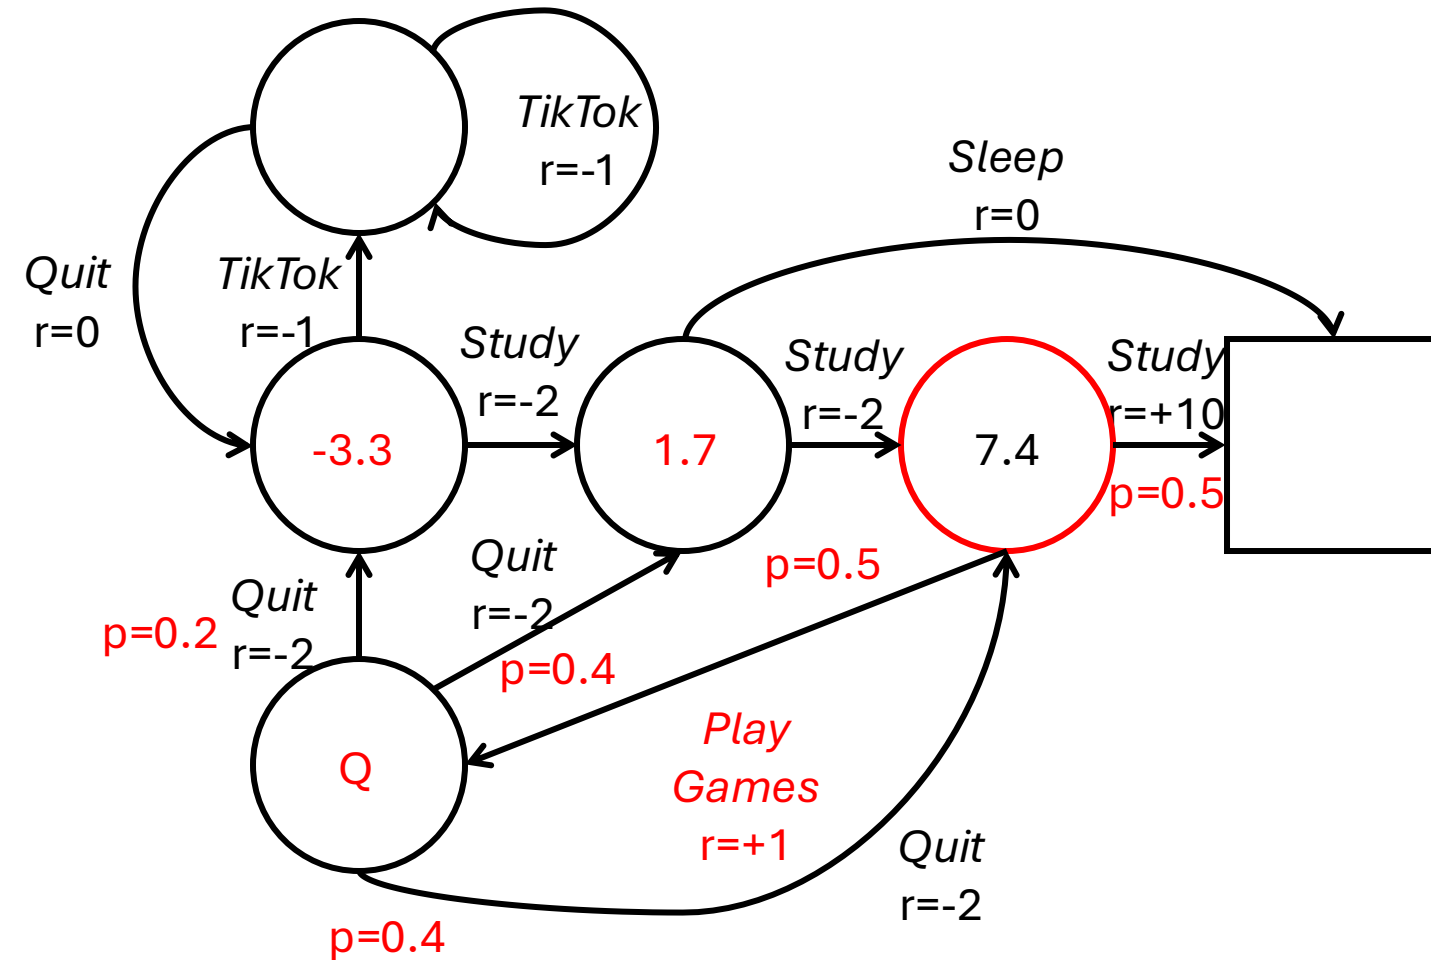
Let's see how we compute V_π for the Class 3 state with $\gamma = 1$

Evaluating a policy with Bellman expectation



$$V_{\pi} = 0.5 * 10 \\ + 0.5 * Q_{\pi}(\text{VG}, \text{Play VG})$$

Evaluating a policy with Bellman expectation



$$\begin{aligned}
 V_{\pi} &= 0.5 * 10 \\
 &+ 0.5 * Q_{\pi}(\text{VG}, \text{Play VG}) \\
 &= 7.4
 \end{aligned}$$

$$\begin{aligned}
 Q_{\pi}(\text{VG}, \text{Play VG}) \\
 &= (1 + 0.2 * -3.3 + 0.4 * 1.7 \\
 &+ 0.4 * 7.4)
 \end{aligned}$$

Closed-form solution

The Bellman expectation equation can be represented as a system of linear equations which results in a closed-form solution:

$$V_{\pi} = R_{\pi} + \gamma T_{\pi} V_{\pi}$$

$$V_{\pi} = (I - \gamma T_{\pi})^{-1} R_{\pi}$$

R_{π} is an $|\mathcal{S}|$ -dimensional vector where j -th entry = $\mathbb{E}[r|s_j, a = \pi(s_j)]$

V_{π} is an $|\mathcal{S}|$ -dimensional vector where j -th entry = $V_{\pi}(s_j)$

T_{π} is an $|\mathcal{S}| \times |\mathcal{S}|$ -dimensional matrix where $(j,k) = p(s_k|s_j, a = \pi(s_j))$

How do we know if a value function is optimal?

The *optimal state-value* function is the max over all policies

$$V^*(s) = \max_{\pi} V_{\pi}(s)$$

And similarly for the *optimal action-value* function...

$$Q^*(s, a) = \max_{\pi} Q_{\pi}(s, a)$$

This represents the best possible performance for a given MDP

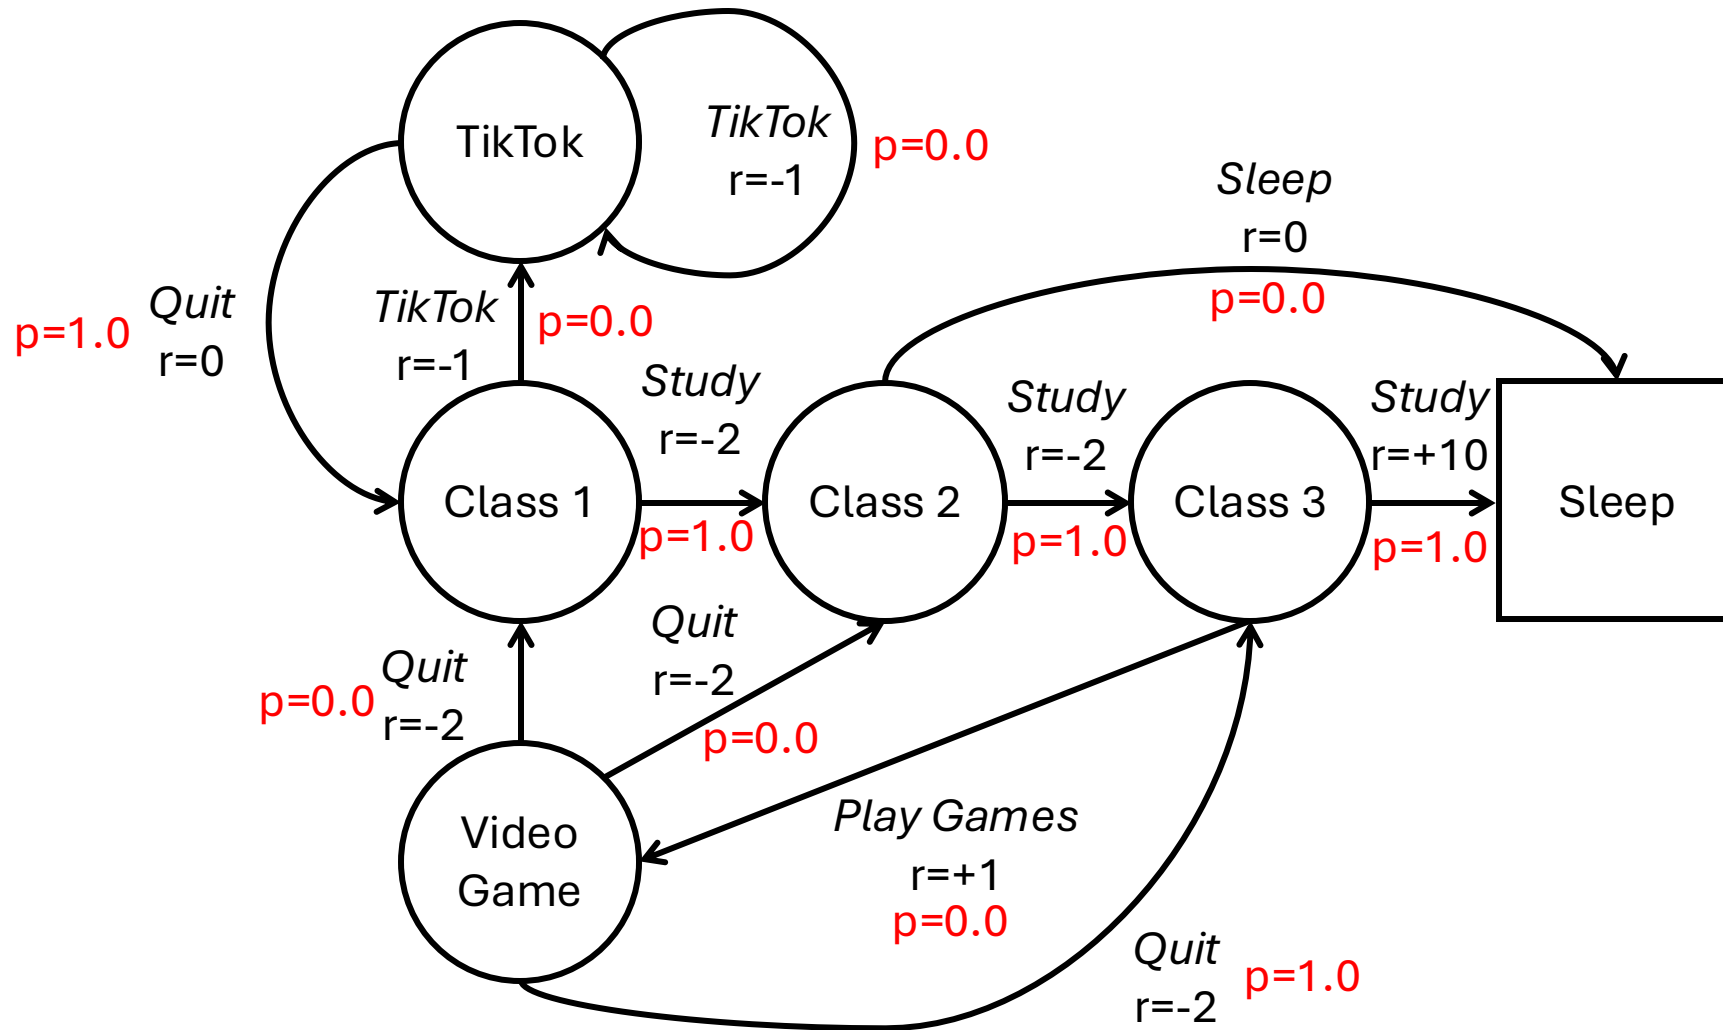
Bellman optimality equations

The value of a state under an optimal policy must equal the expected return for the best action from that state:


$$V^*(s) = \max_{a \in \mathcal{A}} \sum_{s', r} p(s', r | s, a) (r + \gamma V^*(s'))$$

$$Q^*(s, a) = \sum_{s', r} p(s', r | s, a) [r + \gamma \max_{a' \in \mathcal{A}} Q^*(s', a')]$$


The optimal policy for our example



The optimal value function for our example



The optimal value function for our example



Known Models: Planning with Dynamic Programming

Fully known model = planning

When we have complete knowledge of the environment transition function and reward function this becomes a **planning problem**!

We can compute the optimal action-value $Q^*(s, a)$ and state-value $V^*(s)$ functions which allows us to “solve” the MDP

- If we have optimal action-value then we have the optimal policy
- If we have optimal state-value and the environment transition function then we have the optimal policy

No general closed-form solution

Bellman optimality equations are non-linear due to the max operator

This means we cannot find the optimal policy by solving a system of linear equations!

Instead...we use an **iterative approach**

Dynamic Programming

What is dynamic programming?

Dynamic Programming

An optimization method and programming paradigm where the overall problem is broken into simpler sub-problems

It consists of two steps:

1. Solve the sub-problems
2. Combine sub-problem solutions to obtain overall solution

Policy iteration: finding the optimal policy

Idea: use Dynamic Programming to find *optimal policy* for a given MDP

Two-step iterative algorithm. Given a policy π ...

- **Evaluate** the policy

$$V_{\pi}(s) = \mathbb{E}[r_{t+1} + \gamma r_{t+2} + \dots | s_t = s]$$

- **Improve** the policy by acting greedily with respect to V_{π}

$$\pi' = \text{greedy}(V_{\pi})$$

Policy iteration intuition

Starting with a random policy, we evaluate it to find V_π

- Key insight: this policy may not be greedy!
- It might not always choose the action that maximizes the immediate expected return based on V_π

By generating a new policy which **is** greedy with respect to V_π , we...

- Make it a little more “greedy” each update
- Monotonic improvement which provably converges to optimum

Iterative policy evaluation

Earlier we discussed a closed-form solution for evaluation

$$V_{\pi} = (I - \gamma T_{\pi})^{-1} r_{\pi}$$

Unfortunately, this solution has a serious flaw. **What is it?**

Iterative policy evaluation

Earlier we discussed a closed-form solution for evaluation

$$V_{\pi} = (I - \gamma T_{\pi})^{-1} R_{\pi}$$

Unfortunately, this solution has a serious flaw. **What is it?**

It is not computationally tractable!

- Requires us to invert $\mathcal{O}(n^3)$ an $|\mathcal{S}| \times |\mathcal{S}|$ -dimensional matrix
- In large state spaces, we have both compute and memory issues

Iterative policy evaluation


We instead iteratively apply the Bellman equations to convergence

1. Randomly initialize our value function $V_0(s)$ for all $s \in \mathcal{S}$
2. For iteration $k = 1, 2, \dots$

Update $V_{k+1}(s)$ from $V_k(s')$ for all $s \in \mathcal{S}$


If $\max_s |V_{k+1}(s) - V_k(s)| < \epsilon$ then stop

Iterative policy evaluation



$$V_{k+1}(s) = \sum_{a \in \mathcal{A}} \pi(a|s) (r + \gamma \sum_{s' \in \mathcal{S}} p(s'|s, a) V_k(s'))$$

Example: Gridworld



	1	2	3
4	5	6	7
8	9	10	11
12	13	14	

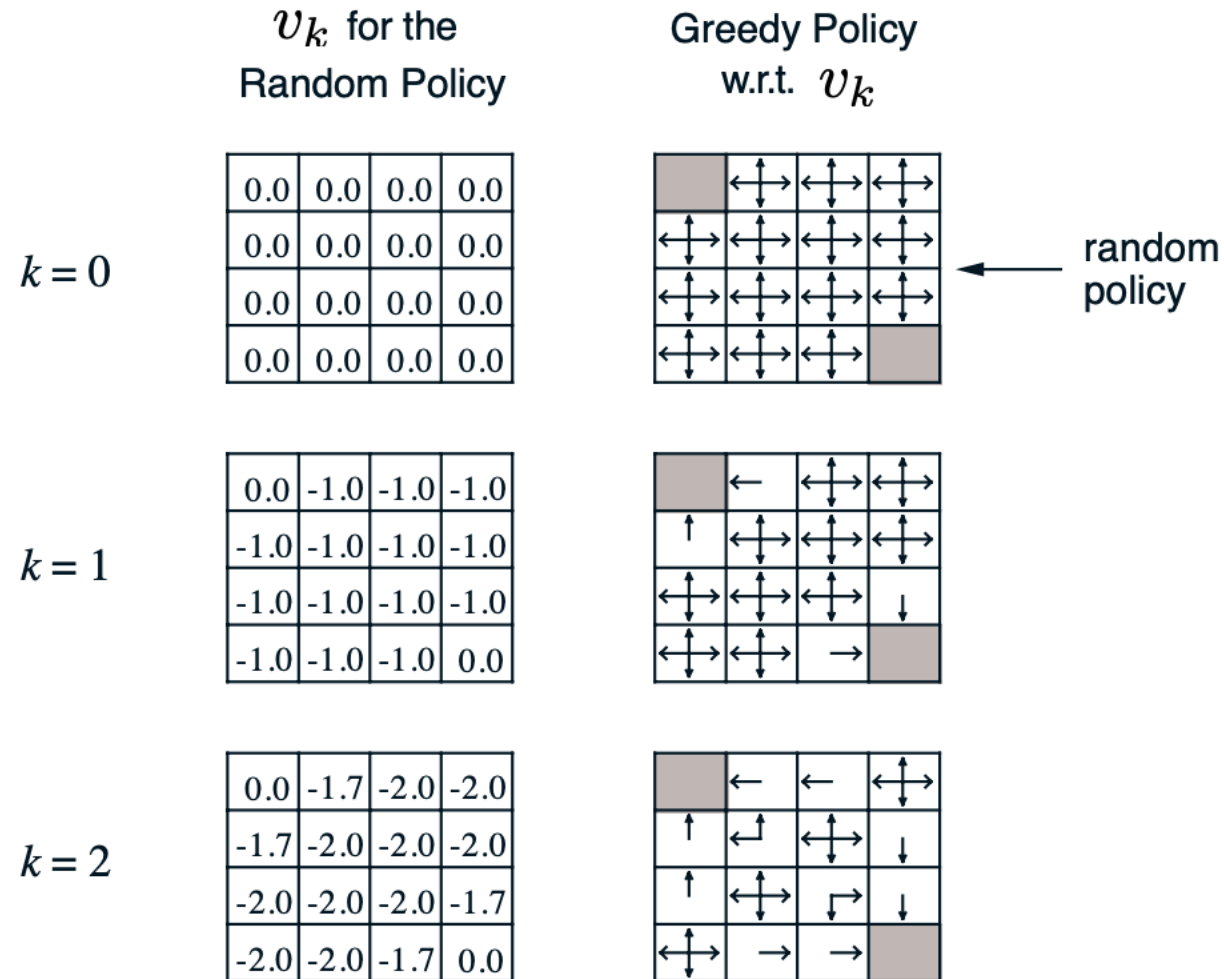
$r = -1$
on all transitions

Terminal states

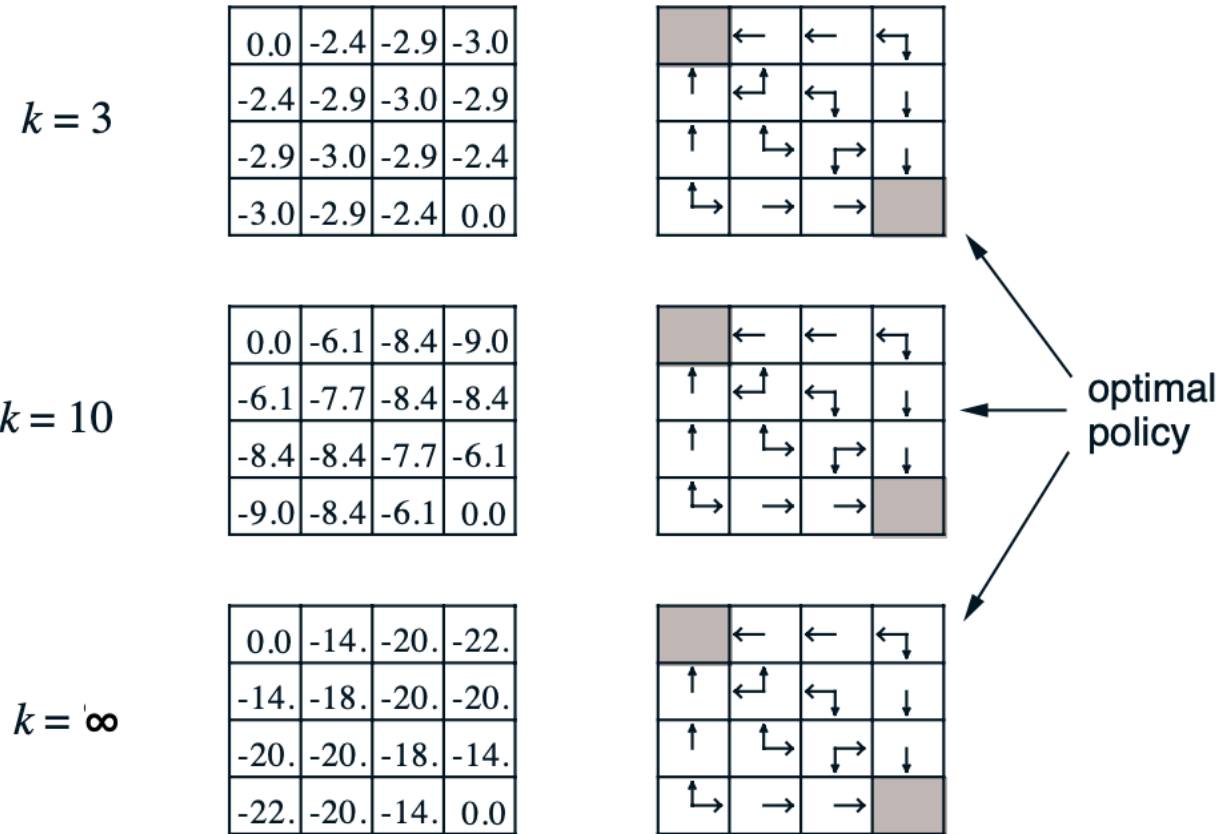
Agent policy is uniformly random

- 25% chance to go north, east, south, or west

Example: Gridworld



Example: Gridworld



Policy improvement

Recall that we know the reward and transition function of the MDP

- This means from state s we know all possible successor states s'

To act greedily, we select actions with the highest expected return

$$\pi'(s) = \operatorname{argmax}_{a \in \mathcal{A}} \sum_{s', r} p(s', r | s, a) (r + \gamma V_{\pi}(s'))$$

Policy improvement

Recall that v^* is the optimal value function of the MDP

- This means $v^*(s) = \max_{a \in \mathcal{A}} \sum_{s', r} P(s', r | s, a) [v^*(s') + R(s, a, s')]$ for states s'

To act greedily with respect to v^* is to choose the action that maximizes the expected return

Guaranteed to converge in finite-horizon MDPs
and discounted infinite-horizon MDPs.

Proof: Sutton and Barto, Chapter 4.2

$$a \in \mathcal{A} \quad \underbrace{\quad}_{s', r}$$

Value iteration: finding the optimal value fn

Idea: rather than computing the value function with respect to π , what if we compute the optimal value function directly?

1. Randomly initialize our value function $V_0(s)$ for all $s \in \mathcal{S}$

2. For iteration $k = 1, 2, \dots$

Update $V_{k+1}(s)$ from $V_k(s')$ for all $s \in \mathcal{S}$

If $\max_s |V_{k+1}(s) - V_k(s)| < \epsilon$ then stop

3. Compute policy from optimal value function V^*

Value iteration: finding the optimal value fn


Idea: rather than computing the value function with respect to π , what if we compute the optimal value function directly?

Unlike iterative policy evaluation, V_k is not with respect to any explicit policy!

Intermediate value functions may not correspond to any policy at all.

3. Compute policy from optimal value function V^*

Value iteration



$$V_{k+1}(s) = \max_{a \in \mathcal{A}} \left(r + \gamma \sum_{s' \in \mathcal{S}} p(s'|s, a) V_k(s') \right)$$

Value iteration

To obtain the resulting policy, act greedily as before

$$\pi(s) = \operatorname{argmax}_{a \in \mathcal{A}} \sum_{s', r} p(s', r | s, a) (r + \gamma V^*(s'))$$

Value iteration

To obtain the resulting policy, act greedily as before

Guaranteed to converge in finite-horizon MDPs and discounted infinite-horizon MDPs.

Proof: Sutton and Barto, Chapter 4.4

Unknown Models: Monte Carlo Learning

The problem with known models

Policy and Value Iteration require fully-known MDPs:

- We must know the reward and transition functions

Is this a problem? Why?

The problem with known models

Policy and Value Iteration require fully-known MDPs:

- We must know the reward and transition functions

Is this a problem? Why?

Yes! In most practical applications we will know **neither**!

- Games, robotics, chatbots, ...

Monte Carlo reinforcement learning

Rather than relying on known environment models, we want to...

- Learn directly from experience (episodes)
- Require no knowledge of transitions/rewards

This is known as **model-free learning**!

Monte Carlo policy evaluation

Given a policy which generates episodes of experience...

$$s_1, a_1, r_2, s_2, a_2, r_3, \dots \sim \pi$$

Previously, the value function is the expected return -- the weighted average of all possible returns that could be obtained from state s :

$$V_{\pi} = \mathbb{E}[G_t | s_t = s]$$

Monte Carlo policy evaluation

Given a policy which generates episodes of experience...

$$s_1, a_1, r_2, s_2, a_2, r_3, \dots \sim \pi$$

Previously, the value function is the expected return -- the weighted average of all possible returns that could be obtained from state s :

$$V_{\pi} = \mathbb{E}[G_t | s_t = s]$$

In Monte Carlo learning we use the **empirical mean** of rewards from experience instead of the **expected return**

- *This is an approximation of the expected return!*

Monte Carlo policy evaluation

Idea: calculate the value of a state as the average of returns observed after visiting that state, using episodes sampled from π

Whenever state s is visited in an episode,

1. Increment visitation counter: $N(s) \leftarrow N(s) + 1$
2. Increment total return: $S(s) \leftarrow S(s) + G_t$

Estimate value by mean return: $V(s) = S(s)/N(s)$

- By the law of large numbers, $V(s) \rightarrow V_\pi(s)$ as $N(s) \rightarrow \infty$

Monte Carlo policy improvement

We can evaluate policies, but how can we **improve** them?

- What has to change about Policy Iteration?

Recall...

$$\pi'(s) = \operatorname{argmax}_{a \in \mathcal{A}} \sum_{s', r} p(s', r | s, a) (r + \gamma V_{\pi}(s'))$$

Monte Carlo policy improvement

We can evaluate policies, but how can we **improve** them?

- What has to change about Policy Iteration?

Recall...

We don't have **successor states**!

$$\pi'(s) = \operatorname{argmax}_{a \in \mathcal{A}} \sum_{\substack{s', r}} p(s', r | s, a) (r + \gamma V_{\pi}(s'))$$

Monte Carlo policy improvement

We can evaluate policies, but how can we **improve** them?

- What has to change about Policy Iteration?

Solution: use the action-value function instead

- No successor states / transition function needed

$$\pi'(s) = \operatorname{argmax}_{a \in \mathcal{A}} Q(s, a)$$

MC policy improvement with action-values

To obtain action-value empirical means instead of state-value...

MC policy improvement with action-values

To obtain action-value empirical means instead of state-value...

Whenever state-action (s, a) is visited in an episode,

1. Increment visitation counter: $N(s, a) \leftarrow N(s, a) + 1$
2. Increment total return: $S(s, a) \leftarrow S(s, a) + G_t$

Estimate value by mean return: $Q(s, a) = S(s, a)/N(s, a)$

MC policy improvement with action-values

There is an important assumption being made, what is it?

Hint: think about step 2

- Increment total return: $S(s, a) \leftarrow S(s, a) + G_t$

MC policy improvement with action-values

There is an important assumption being made, what is it?

Hint: think about step 2

- Increment total return: $S(s, a) \leftarrow S(s, a) + G_t$

To compute G_t we must have complete episodes!

- This means all episodes must terminate

Exploration-exploitation dilemma revisited

Since we're computing action-values using empirical means, we have to actually *try* sub-optimal actions to learn their values

We want to learn action-values for an optimal policy...but we need to act suboptimally to explore all actions

Exploration-exploitation dilemma revisited

Since we're computing action-values using empirical means, we have to actually *try* sub-optimal actions to learn their values

We want to learn action-values for an optimal policy...but we need to act sub-optimally to explore all actions

Easy solution: ϵ -greedy

Recap: ϵ -greedy

$$\pi(a|s) = \begin{cases} \epsilon/m + 1 - \epsilon, & \text{if } a^* = \operatorname{argmax}_{a \in \mathcal{A}} Q(s, a) \\ \epsilon/m, & \text{otherwise} \end{cases}$$



For m actions

Recap: ϵ -greedy

$$\pi(a|s) = \begin{cases} \epsilon/m + 1 - \epsilon, & \text{if } a^* = \underset{a \in \mathcal{A}}{\operatorname{argmax}} Q(s, a) \\ \epsilon/m, & \text{otherwise} \end{cases}$$

Monte Carlo policy improvement with an ϵ -greedy policy is guaranteed to converge to the optimal action-value function

- *Caveat: with infinite exploration*
- Proof: Sutton / Barto, Chapter 5.4

Monte Carlo on-policy learning

Two-step iterative algorithm. Randomly initialize policy π ...

- **Evaluate** the policy with sampled episodes

$Q_{\pi}(s, a)$ approximated with empirical means

- **Improve** the policy by acting ϵ -greedily with respect to V_{π}

$$\pi' = \epsilon\text{-greedy } Q(s, a)$$

Note: consider decaying ϵ to converge to an optimal policy

What is “on-policy” learning?

On-policy learning:

- “Learn on the job”
- The policy learns from its own experience
- Improve policy π from episodes sampled from π

Off-policy learning:

- “Look over someone’s shoulder”
- The policy learns from another policy’s experience
- Improve policy π from episodes sampled from β

Take-aways

- We evaluate policy quality by decomposing value functions with the Bellman equations
- In fully known MDPs (reward + transition), find the optimal policy using Policy Iteration and Value Iteration
- When we don't know the full model, use Monte Carlo methods

Next time...

Temporal Difference learning (non-complete episodes?)

- This forms the basis of modern reinforcement learning

Introduction to deep reinforcement learning (and DQN)

Interactive development session with Pytorch / TorchRL?