

RCAC Tutorial

Guven Gergerli

RCAC Scholar

- Groups of computer clusters for research at Purdue
- Job scheduling or remote desktop options
- There are different Community Clusters but we will use Scholar
 - classroom learning about high performance computing (HPC).
- For anything that is not mentioned, there is a detailed guide on how to use Scholar.

The screenshot displays the RCAC Scholar website. At the top, the Purdue University logo and the Rosen Center for Advanced Computing name are visible. A navigation bar contains links for News, Support Hub, Engagement, Account, Compute, Storage, Anvil, RSE, Purchase, Services, and About. Below this, a breadcrumb trail reads: Home / Knowledge Base / Scholar User Guide / Overview of Scholar. The main content area features a search bar and a sidebar with a list of links: Bell User Guide, Gilbreth User Guide, Weber User Guide, Scholar User Guide (expanded), Overview of Scholar (selected), Accounts, File Storage and Transfer, Gateway (Open OnDemand), Software, Compiling Source Code, Running Jobs, Using Jupyter Hub on Scholar, Frequently Asked Questions, and Datasets. The 'Overview of Scholar' section contains the following text:

Overview of Scholar

Scholar is a small computer cluster, suitable for classroom learning about high performance computing (HPC). It consists of 6 interactive login servers and 16 batch worker nodes.

It can be accessed as a typical cluster, with a job scheduler distributing batch jobs onto its worker nodes, or as an interactive resource, with software packages available through a desktop-like environment on its login servers.

If you have a class that you think will benefit from the use of Scholar, you can schedule it for your class through our [Class Account Request](#) page. You only need to register your class itself. All students who register for the class will automatically get login privileges to the Scholar cluster.


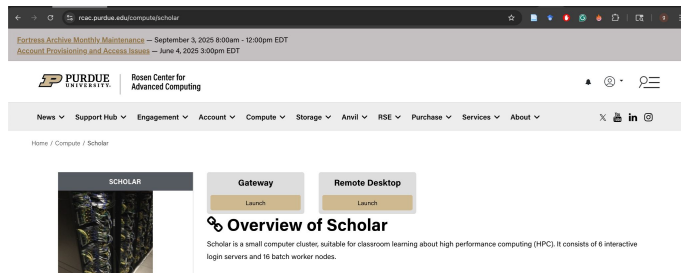
As a batch resource, the cluster has access to typical HPC software packages and tool chains; as an interactive resource, Scholar provides a Linux remote desktop, or a Jupyter notebook server, or an R Studio server. Jupyter and R Studio can be used by students without any reliance on Linux knowledge or experience.

How to use Scholar for GPU training

1. Gateway

- Many functionalities with browser access
- This is what you will see:

rcac.purdue.edu/compute/scholar



Remote Desktop



OPEN

OnDemand

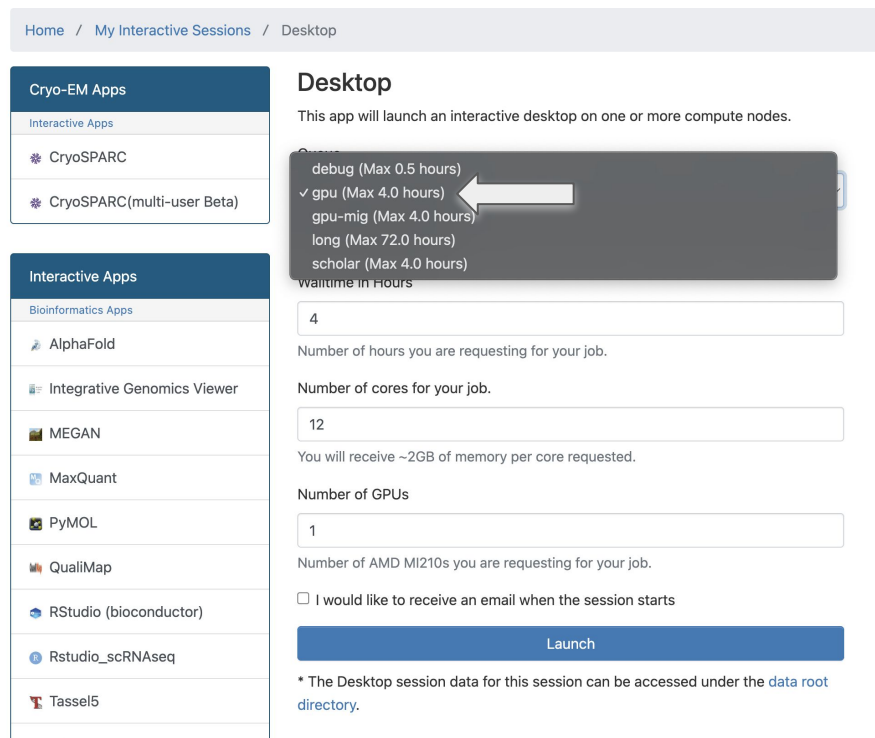
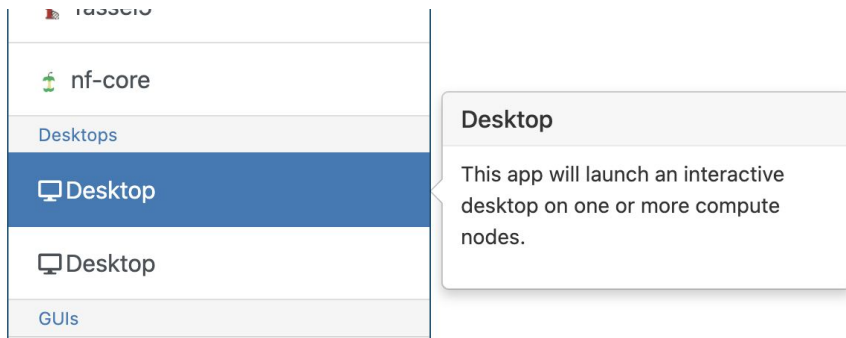
OnDemand provides an integrated, single access point for all of your HPC resources.

Message of the Day

How to use Scholar for GPU training

2. Remote Desktop

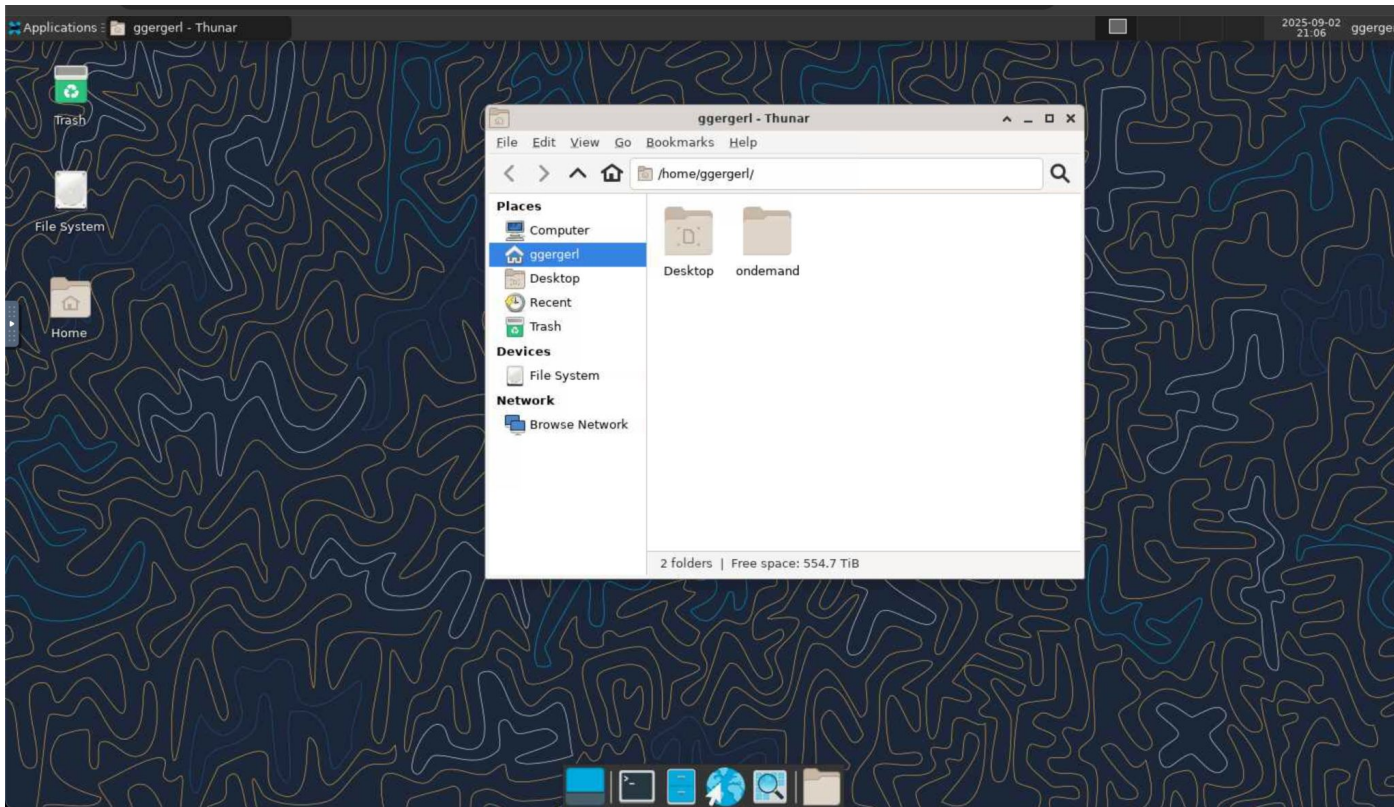
- Best to initially see what is going on inside the cluster



- This configuration if you would like to train inside the remote desktop

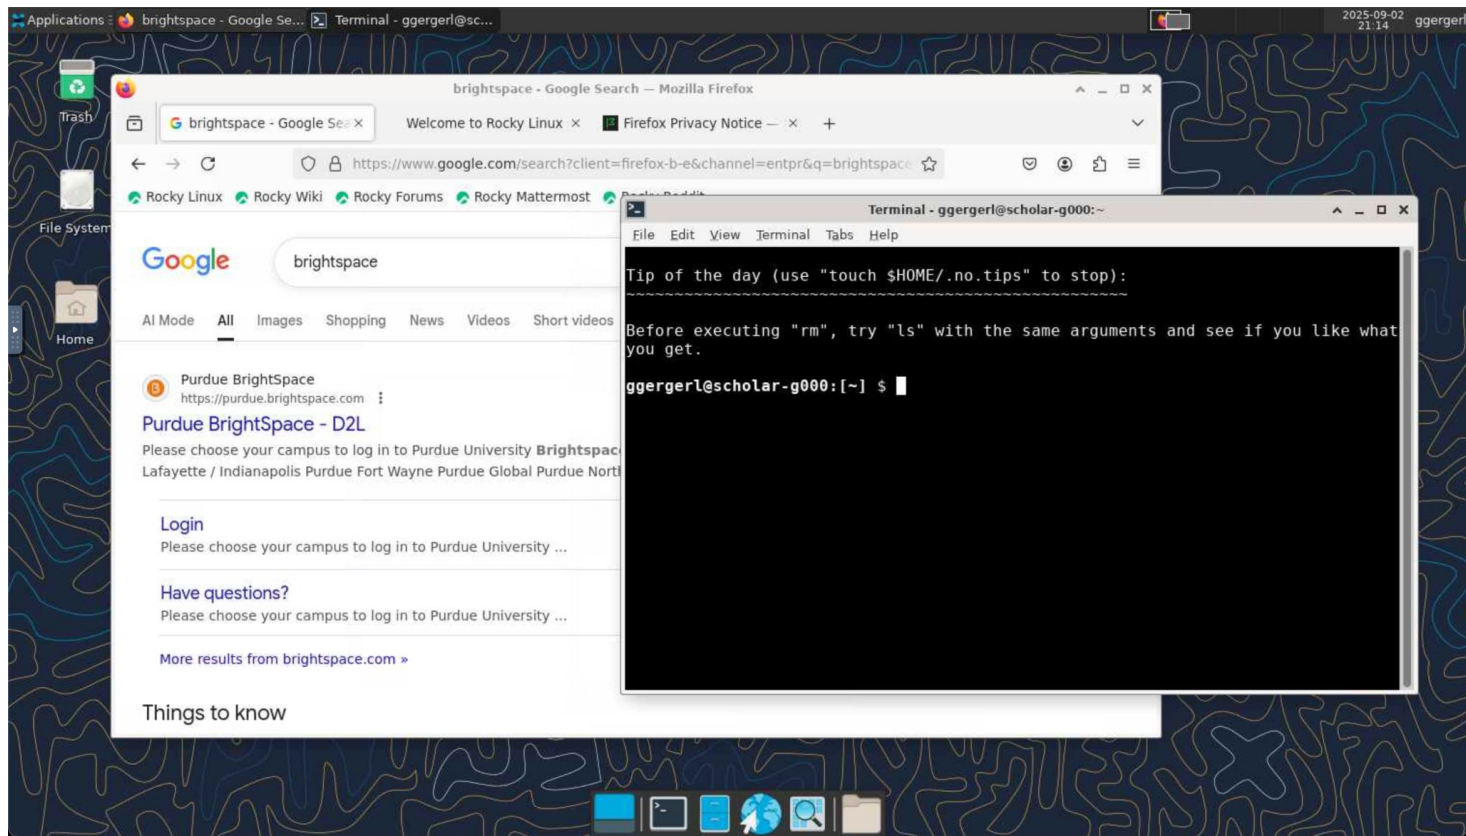
How to use Scholar for GPU training

2. Remote Desktop



How to use Scholar for GPU training


2. Remote Desktop



How to use Scholar for GPU training

3. Shell

- When you need to access the terminal through the browser (without having to start remote desktop)



```
*****
***** Use of Purdue BoilerKey or SSH keys is Required *****
*****

Last login: Tue Sep  2 20:58:17 2025 from gateway.scholar.rcac.purdue.edu

Tip of the day (use "touch $HOME/.no.tips" to stop):
~~~~~

Vimdiff is a vi based tool to merge files.

ggergerl@scholar-fe00:[~] $
ggergerl@scholar-fe00:[~] $ ls
Desktop  ondemand
ggergerl@scholar-fe00:[~] $ pwd
/home/ggergerl
ggergerl@scholar-fe00:[~] $ █
```

How to use Scholar for GPU training

4. Running jobs

- Clusters can be initiated with a SLURM job
 - SLURM job: a configuration for clusters to run a script with dedicated hardware
- Remote Desktop is also a job but with a desktop application
- For more detail please check RCAC tutorial as well ->

Search knowledge base...

Submit

Expand Topics

Basics of SLURM Jobs

The *Simple Linux Utility for Resource Management (SLURM)* is a system providing job scheduling and job management on compute clusters. With SLURM, a user requests resources and submits a job to a queue. The system will then take jobs from queues, allocate the necessary nodes, and execute them.

Do NOT run large, long, multi-threaded, parallel, or CPU-intensive jobs on a front-end login host. All users share the front-end hosts, and running anything but the smallest test job will negatively impact everyone's ability to use Scholar. Always use SLURM to submit your work as a job.

Submitting a Job

The main steps to submitting a job are:

- [Create job submission script](#)
- [Submit job script](#)
- [Monitor job status](#)
- [Check output](#)

Follow the links below for information on these steps, and other basic information about jobs. A number of [example SLURM jobs](#) are also available.

- [Queues](#)
- [Job Submission Script](#)

> Bell User Guide

> Gilbreth User Guide

> Weber User Guide

> Scholar User Guide

> Overview of Scholar

> Accounts

> File Storage and Transfer

> Gateway (Open OnDemand)

> Software

> Compiling Source Code

> Running Jobs

> Basics of SLURM Jobs

> Queues

> Job Submission Script

> Submitting a Job

> Job Dependencies

> Holding a Job

> Checking Job Status

> Checking Job Output

> Canceling a Job

> Example Jobs

> Using Jupyter Hub on Scholar

> Frequently Asked Questions

How to use Scholar for GPU training

4. Running jobs

- You can run jobs with a job submission file, this is what it looks like:

An example from RCAC

```
#!/bin/bash
# FILENAME: myjobsubmissionfile

# Loads Matlab and sets the application up
module load matlab

# Change to the directory from which you originally submitted this job.
cd $SLURM_SUBMIT_DIR

# Runs a Matlab script named 'myscript'
matlab -nodisplay -singleCompThread -r myscript
```

- You can specify within scholars own limitations.
- The files should be ``.sub`` and ran through ``sbatch file_name.sub``

One of my customized job submission file
(for another task on a different cluster)

```
1  #!/bin/bash
2  #SBATCH --job-name=ppo-burr
3  #SBATCH --output=errout/rl-%J.out
4  #SBATCH --error=errout/rl-%J.err
5  #SBATCH --cpus-per-task=12
6  #SBATCH --time=30:00:00
7  #SBATCH --account=jocamp
8  #SBATCH --partition=a100-40gb
9  #SBATCH --nodes=1
10 #SBATCH --mem=240G
11 #SBATCH --gpus-per-node=1
12
13
14 # Load modules
15 module load cuda
16 module load conda
17
18
19 export PYTHONPATH=$PYTHONPATH:~/GitHub/multiagent-tom:~/GitHub/ToMOvercooked_ai
20
21 # Activate conda environment
22 source $(conda info --base)/etc/profile.d/conda.sh
23 conda activate overcooked_ai_lucy
24
25
26 python ~/GitHub/multiagent-tom/src/burrito_torchrl/ppo.py --config new_config.yaml
27
```

How to use Scholar for GPU training

4. Running jobs

- A job that is on the process can be seen in the active jobs in gateway

Scholar

Files

Jobs

Clusters

Interactive Apps

Cryo-EM Apps

Krannert Apps

My Interactive Sessions

Help

Logged in as ggergerl

Log Out

Active Jobs

Active Jobs

Job Composer

Your Jobs

All Clusters

Active Jobs

Show 50 entries

Filter:

ID	Name	User	Account	Time Used	Queue	Status	Cluster	Actions	
>	382420	OnDemand/Desktop	ggergerl	gpu	00:22:41	scholar-gpu	Running	Scholar	

Showing 1 to 1 of 1 entries

Previous1Next

How to use Scholar for GPU training

Bonus: VSCode

- You can also use VSCode for directly playing around with your code inside the cluster
- For this you will need to configure an SSH Host
 - For this you should check the tutorial on RCAC

➤ Bell User Guide

➤ Gilbreth User Guide

➤ Weber User Guide

▼ Scholar User Guide

➤ Overview of Scholar

▼ Accounts

▼ Logging In

➤ Purdue Login

➤ Passwords

➤ SSH Client Software

➤ **SSH Keys**

➤ ThinLinc

➤ SSH X11 Forwarding

➤ Purchasing Nodes

➤ File Storage and Transfer

➤ Gateway (Open OnDemand)

➤ Software

➤ Compiling Source Code

➤ Running Jobs

➤ Using Jupyter Hub on Scholar

Search knowledge base...

Submit

Expand Topics

SSH Keys

General overview

To connect to Scholar using SSH keys, you must follow three high-level steps:

1. Generate a key pair consisting of a private and a public key on your local machine.
2. Copy the public key to the cluster and append it to `~/.ssh/authorized_keys` file in your account.
3. Test if you can ssh from your local computer to the cluster without using your Purdue password.

Detailed steps for different operating systems and specific SSH client softwares are give below.

Mac and Linux:

1. Run `ssh-keygen` in a terminal on your local machine. You may supply a filename and a passphrase for protecting your private key, but it is not mandatory. To accept the default settings, press Enter without specifying a filename.
Note: If you do not protect your private key with a passphrase, anyone with access to your computer could SSH to your account on Scholar.
2. By default, the key files will be stored in `~/.ssh/id_rsa` and `~/.ssh/id_rsa.pub` on your local machine.

Select an option to open a Remote Window

Connect to Host...	Remote-SSH
Connect Current Window to Host...	
Connect to Tunnel...	Remote-Tunnels
Open Container Configuration File	Dev Containers

Select configured SSH host or enter user@host

gilbreth.rcac.purdue.edu

+ Add New SSH Host...

Configure SSH Hosts...

Enter SSH Connection Command

E.g. `ssh hello@microsoft.com -A`

Press 'Enter' to confirm your input or 'Escape' to cancel